

Can a small AI model summarize job postings for career advisers?

Shreya Verma · Chih-yu Li · William Sun | Pharma · Finance · Tech job postings

95820 Applications of NLX and LLM · Carnegie Mellon University · October 2026

The question


A university career office wants short, **source-linked summaries** of job postings (title, employer, location, work arrangement, requirements) that an adviser checks before students see them. Not hiring decisions.

The data: 650 public job postings

Sector	Records	Source URLs	Employers
Pharma	200	200	8
Finance	200	175	27
Tech	250	71	68

Collected Sept 18–20, 2026 from public job-board APIs (Greenhouse, Lever), NYC Open Data and employer career pages. Tech rows are often fragments of one posting, so more rows do not mean more jobs.


How the API works



Everything runs on a laptop: no posting leaves the machine. We tested six designs on 50 postings each (300 requests), then attacked the API with 30 safety probes.

Finding 1: tool works, answers don’t

The lookup tool ran 50/50 times. But only 17/50 answers got all five source facts right.



Five facts: posting ID, URL, title, employer, location. One run per design.

Finding 2: formatting hid answers

The model wrapped its JSON in Markdown fences, so the strict checker rejected **all 50** structured answers. Stripping the fences afterwards gave **43/50** valid, **31/50** source-correct: a diagnosis, not a new result.


A checker that refuses anything unverified (C4) blocked all 50 answers. Perfectly safe, completely useless.

Same lesson in all three sectors

Sector (AS01 test set)	Phi valid JSON	Larger hosted model
Pharma (25)	23/25	Gemini Flash: 25/25
Finance (28)	12/28	gpt-oss-120b: 28/28
Tech (25)	7/25	GPT-5-mini: 24/25


Bigger models fixed the format; skill and qualification lists stayed hard. Schemas differ, so this is not a ranking; Finance/Tech figures from teammate memos.

Finding 3: safety has a service cost



Keyword guardrail blocked 15/20 attacks but also **1 in 10** normal requests (above our 5% limit). It agreed with human judgment on **24/30** probes. Misses were paraphrases.

Finding 4: review time is the cost



From single timings in an earlier prototype; if review took 249 s, the saving disappears.

Recommendation

- Copy facts already in the source table directly: no AI needed.
- Use the AI only as a supervised research tool, with an adviser checking every answer.
- Never for ranking applicants or eligibility decisions.
- Next: enforce JSON at decoding time, test on new employers.